

Chapter 18: B-trees

Algorithms and Data Structures

(based on book “Introduction to Algorithms” of Cormen et al.)

Vincent Van Schependom

KU Leuven Campus Kulak Kortrijk

Academic year 2024–2025

Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

Definition of a B-tree

- ▶ Rooted tree with root $T.root$

Definition of a B-tree

- ▶ Rooted tree with root $T.root$
- ▶ Self-balancing, just like red-black trees.

Definition of a B-tree

- ▶ Rooted tree with root $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height h .

Definition of a B-tree

- ▶ Rooted tree with root $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height h .
- ▶ Every node x has $x.n$ keys, stored in monotonically increasing order:

$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$

Definition of a B-tree

- ▶ Rooted tree with root $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height h .
- ▶ Every node x has $x.n$ keys, stored in monotonically increasing order:

$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$

- ▶ Each internal node contains $x.n + 1$ pointers to its children:

$$x.c_1, x.c_2, \dots, x.c_{x.n+1}$$

Definition of a B-tree

- ▶ Rooted tree with root $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height h .
- ▶ Every node x has $x.n$ keys, stored in monotonically increasing order:


$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$

- ▶ Each internal node contains $x.n + 1$ pointers to its children:

$$x.C_1, x.C_2, \dots, x.C_{x.n+1}$$

- ▶ Main concept \rightarrow remember well

Example of a B-tree




1 node,
1000 keys

1001 nodes,
1,001,000 keys

1,002,001 nodes,
1,002,001,000 keys

- Every node x has $x.n = 1000$ keys $x.key_i$.

Example of a B-tree



1 node,
1000 keys

1001 nodes,
1,001,000 keys

1,002,001 nodes,
1,002,001,000 keys

- ▶ Every node x has $x.n = 1000$ keys $x.key_i$.
- ▶ Each internal node x has $x.n + 1 = 1001$ pointers to its children $x.c_i$.

Key ordering

- ▶ The keys $x.key_i$ separate the ranges of keys stored in each subtree.

Key ordering

- ▶ The keys $x.key_i$ separate the ranges of keys stored in each subtree.
 - Let k_i be any key stored in the subtree with root $x.c_i$

Key ordering

- ▶ The keys $x.key_i$ separate the ranges of keys stored in each subtree.
 - Let k_i be any key stored in the subtree with root $x.c_i$


$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$

Key ordering

- ▶ The keys $x.key_i$ separate the ranges of keys stored in each subtree.
 - Let k_i be any key stored in the subtree with root $x.c_i$

$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$

- ▶ Example:



Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least?

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least? t children.

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least? t children.
 - If the tree is nonempty, the root must have at least one key.

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least? t children.
 - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least? t children.
 - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
 - Every node may contain at most $2t - 1$ keys.

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least? t children.
 - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
 - Every node may contain at most $2t - 1$ keys.
 - So, how many children can an internal node have at most?

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least? t children.
 - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
 - Every node may contain at most $2t - 1$ keys.
 - So, how many children can an internal node have at most? $2t$ children.

Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer $t \geq 2$, called the *minimum degree*
- ▶ Lower bound:
 - Every node other than the root must have at least $t - 1$ keys.
 - So, how many children does an internal node need at least? t children.
 - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
 - Every node may contain at most $2t - 1$ keys.
 - So, how many children can an internal node have at most? $2t$ children.
 - We say that a node is *full* if it contains exactly $2t - 1$ keys.

Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:

Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
 - Fast.

Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
 - Fast.
 - Expensive and limited in capacity.

Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
 - Fast.
 - Expensive and limited in capacity.
- ▶ Secondary memory:

Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
 - Fast.
 - Expensive and limited in capacity.
- ▶ Secondary memory:
 - Cheaper and much larger.

Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
 - Fast.
 - Expensive and limited in capacity.
- ▶ Secondary memory:
 - Cheaper and much larger.
 - We need such capacity for storing large amounts of data, like in databases.


Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory


- ▶ Primary memory:
 - Fast.
 - Expensive and limited in capacity.
- ▶ Secondary memory:
 - Cheaper and much larger.
 - We need such capacity for storing large amounts of data, like in databases.
 - E.g. SSD's and HDD's.

Secondary memory is slow – Example: HDD



- ▶ Because of the moving mechanical parts, HDD's are slow.
- ▶ SSD's are not mechanical, but still significantly slower than RAM.

Secondary memory is slow – Example: HDD



- ▶ Because of the moving mechanical parts, HDD's are slow.
- ▶ SSD's are not mechanical, but still significantly slower than RAM.
→ latency!

Combating latency

- ▶ Access not just one item, but several at a time.

Combating latency

- ▶ Access not just one item, but several at a time.
 - Information is divided into a number of equal-sized *blocks*.

Combating latency

- ▶ Access not just one item, but several at a time.
 - Information is divided into a number of equal-sized *blocks*.
 - B-trees are a very common datastructure for handling this.

Combating latency

- ▶ Access not just one item, but several at a time.
 - Information is divided into a number of equal-sized *blocks*.
 - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.

Combating latency

- ▶ Access not just one item, but several at a time.
 - Information is divided into a number of equal-sized *blocks*.
 - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory \leftrightarrow operations on a B-tree

Combating latency

- ▶ Access not just one item, but several at a time.
 - Information is divided into a number of equal-sized *blocks*.
 - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory \leftrightarrow operations on a B-tree
 - The number of blocks read or written provides a good approximation of the total time spent accessing the disk drive.

Combating latency

- ▶ Access not just one item, but several at a time.
 - Information is divided into a number of equal-sized *blocks*.
 - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory \leftrightarrow operations on a B-tree
 - The number of blocks read or written provides a good approximation of the total time spent accessing the disk drive.
 - Only downward paths.

Combating latency

- ▶ Access not just one item, but several at a time.
 - Information is divided into a number of equal-sized *blocks*.
 - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory \leftrightarrow operations on a B-tree
 - The number of blocks read or written provides a good approximation of the total time spent accessing the disk drive.
 - Only downward paths.
 - So we want the height to be as small as possible.

Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

Height of a B-tree

Theorem

If $n \geq 1$, then for any n -key B-tree T of height h and minimum degree $t \geq 2$


$$h \leq \log_t \left(\frac{n+1}{2} \right)$$

Height of a B-tree

Theorem

If $n \geq 1$, then for any n -key B-tree T of height h and minimum degree $t \geq 2$

$$h \leq \log_t \left(\frac{n+1}{2} \right)$$




Height of a B-tree

Theorem

If $n \geq 1$, then for any n -key B-tree T of height h and minimum degree $t \geq 2$


$$h \leq \log_t \left(\frac{n+1}{2} \right)$$



depth	0	1	2	3	...	h
number of nodes	1	2	$2t$	$2t^2$...	$2t^{h-1}$

Height of a B-tree

Proof.




depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
\vdots	\vdots
h	$2t^{h-1}$

Height of a B-tree

Proof. The number of keys n satisfies

$$n \geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}}$$




depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
\vdots	\vdots
h	$2t^{h-1}$

Height of a B-tree

Proof. The number of keys n satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left(\frac{t^h - 1}{t - 1} \right)
 \end{aligned}$$




depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
\vdots	\vdots
h	$2t^{h-1}$

Height of a B-tree

Proof. The number of keys n satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left(\frac{t^h - 1}{t-1} \right) \\
 &= 2t^h - 1
 \end{aligned}$$



depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
\vdots	\vdots
h	$2t^{h-1}$


Height of a B-tree

Proof. The number of keys n satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left(\frac{t^h - 1}{t - 1} \right) \\
 &= 2t^h - 1
 \end{aligned}$$

Thus

$$t^h \leq \frac{n+1}{2}$$



depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
\vdots	\vdots
h	$2t^{h-1}$


Height of a B-tree

Proof. The number of keys n satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2^{t^{i-1}}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left(\frac{t^h - 1}{t - 1} \right) \\
 &= 2t^h - 1
 \end{aligned}$$

Thus

$$t^h \leq \frac{n+1}{2} \iff h \leq \log_t \left(\frac{n+1}{2} \right)$$



depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
\vdots	\vdots
h	$2t^{h-1}$

Comparison with red-black trees

- ▶ Height of the tree grows as $O(\lg n)$ in both cases.

Comparison with red-black trees

- ▶ Height of the tree grows as $O(\lg n)$ in both cases.
 - Constant t is left out in asymptotic notation.

Comparison with red-black trees

- ▶ Height of the tree grows as $O(\lg n)$ in both cases.
 - Constant t is left out in asymptotic notation.
 - The base logarithm for B-trees is much larger: t vs 2.

Comparison with red-black trees

- ▶ Height of the tree grows as $O(\lg n)$ in both cases.
 - Constant t is left out in asymptotic notation.
 - The base logarithm for B-trees is much larger: t vs 2.
- ▶ Factor of $\approx \lg t$ saved over red-black trees.

Comparison with red-black trees

- ▶ Height of the tree grows as $O(\lg n)$ in both cases.
 - Constant t is left out in asymptotic notation.
 - The base logarithm for B-trees is much larger: t vs 2.
- ▶ Factor of $\approx \lg t$ saved over red-black trees.
- ▶ B-trees are limited in depth but contain an enourmous amount of keys.

Comparison with red-black trees

- ▶ Height of the tree grows as $O(\lg n)$ in both cases.
 - Constant t is left out in asymptotic notation.
 - The base logarithm for B-trees is much larger: t vs 2.
- ▶ Factor of $\approx \lg t$ saved over red-black trees.
- ▶ B-trees are limited in depth but contain an enourmous amount of keys.
 - Few memory operations, so minimal memory latency.

Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.

Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.

Searching a B-tree


- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.
 - The amount of *ways* is equal to the number of children.

Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.
 - The amount of *ways* is equal to the number of children.
 - At each internal node x , the search makes an $(x.n + 1)$ -way branching decision.

Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.
 - The amount of *ways* is equal to the number of children.
 - At each internal node x , the search makes an $(x.n + 1)$ -way branching decision.
- ▶ Example: searching for the letter S using lexicographic order.



Searching a B-tree: pseudocode

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

Searching a B-tree: pseudocode

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

► x is a pointer to the root node of a subtree

Searching a B-tree: pseudocode

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- ▶ x is a pointer to the root node of a subtree
- ▶ k is the key to be searched

Searching a B-tree: pseudocode

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- ▶ x is a pointer to the root node of a subtree
- ▶ k is the key to be searched
- ▶ If k is in the B-tree, B-TREE-SEARCH returns the node-index pair (y, i) such that $y.key_i = k$. Otherwise, the procedure returns NIL.

Searching a B-tree: memory access analysis

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.

Searching a B-tree: memory access analysis

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.
- ▶ Thus, B-TREE-SEARCH executes $O(h) = O(\log_t n)$ memory accesses.

Searching a B-tree: memory access analysis

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.
- ▶ Thus, B-TREE-SEARCH executes $O(h) = O(\log_t n)$ memory accesses.
 - Recall that $h \leq \log_t \left(\frac{n+1}{2}\right)$.

Searching a B-tree: memory access analysis

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.
- ▶ Thus, B-TREE-SEARCH executes $O(h) = O(\log_t n)$ memory accesses.
 - Recall that $h \leq \log_t \left(\frac{n+1}{2}\right)$.
 - Constants are left out in O -notation.

Searching a B-tree: runtime analysis

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- For every node x , $x.n < 2t$, so the while loop takes $O(t)$ time for each encountered node.

Searching a B-tree: runtime analysis

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- ▶ For every node x , $x.n < 2t$, so the while loop takes $O(t)$ time for each encountered node.
- ▶ The total runtime is $O(th) = O(t \log_t n)$.

Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.

Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.
- ▶ We cannot simply create a new leaf node and insert it, because the resulting tree would fail to be a valid B-tree.


Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.
- ▶ We cannot simply create a new leaf node and insert it, because the resulting tree would fail to be a valid B-tree.
- ▶ So, we need to insert the key into an *existing* leaf node.

Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.
- ▶ We cannot simply create a new leaf node and insert it, because the resulting tree would fail to be a valid B-tree.
- ▶ So, we need to insert the key into an *existing* leaf node.
- ▶ What if a node y is *full*, i.e. what if $y.n = 2t - 1$?

Splitting a full node y with $y.n = 2t - 1$



B-TREE-SPLIT-CHILD(x, i)

- ▶ x is a non-full node and $y = x.c_i$ is a full child of x .
- ▶ Split y about its median key S and move S up into y 's parent node x .
- ▶ Every key $y.key_i$ that is greater than the median S , is placed in a new node z , which is a new child of x .

Inserting a key (continued)

Complexity of $\text{B-TREE-INSERT}(T, k)$?

Inserting a key (continued)

Complexity of B-TREE-INSERT(T, k)?

- ▶ Again $O(h) = O(\log_t n)$ disk accesses.

Inserting a key (continued)

Complexity of B-TREE-INSERT(T, k)?

- ▶ Again $O(h) = O(\log_t n)$ disk accesses.
- ▶ Again $O(th) = O(t \log_t n)$ time required.

Questions?