

# Algorithms and Data Structures

Chapter 18: B-trees

(based on book “Introduction to Algorithms” of Cormen et al.)

Vincent Van Schependom

KU Leuven Campus Kulak Kortrijk

Academic year 2024–2025

# Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

# Outline

- 1 Definition of a B-tree
- 2 Use-cases
- 3 Height of a B-tree
- 4 Operations on B-trees

## Definition of a B-tree

- ▶ Rooted tree with root  $T.root$

## Definition of a B-tree

- ▶ Rooted tree with root  $T.root$
- ▶ Self-balancing, just like red-black trees.

## Definition of a B-tree

- ▶ Rooted tree with root  $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height  $h$ .

## Definition of a B-tree

- ▶ Rooted tree with root  $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height  $h$ .
- ▶ Every node  $x$  has  $x.n$  keys, stored in monotonically increasing order:

$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$

## Definition of a B-tree

- ▶ Rooted tree with root  $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height  $h$ .
- ▶ Every node  $x$  has  $x.n$  keys, stored in monotonically increasing order:

$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$

- ▶ Each internal node contains  $x.n + 1$  pointers to its children:

$$x.c_1, x.c_2, \dots, x.c_{x.n+1}$$



## Definition of a B-tree

- ▶ Rooted tree with root  $T.root$
- ▶ Self-balancing, just like red-black trees.
- ▶ All leaves have the same depth, i.e. the tree's height  $h$ .
- ▶ Every node  $x$  has  $x.n$  keys, stored in monotonically increasing order:


$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$

- ▶ Each internal node contains  $x.n + 1$  pointers to its children:

$$x.C_1, x.C_2, \dots, x.C_{x.n+1}$$

- ▶ Main concept  $\rightarrow$  remember well

## Example of a B-tree




1 node,  
1000 keys

1001 nodes,  
1,001,000 keys

1,002,001 nodes,  
1,002,001,000 keys

- Every node  $x$  has  $x.n = 1000$  keys  $x.key_i$ .

## Example of a B-tree



1 node,  
1000 keys

1001 nodes,  
1,001,000 keys

1,002,001 nodes,  
1,002,001,000 keys

- ▶ Every node  $x$  has  $x.n = 1000$  keys  $x.key_i$ .
- ▶ Each internal node  $x$  has  $x.n + 1 = 1001$  pointers to its children  $x.c_i$ .

## Key ordering

- ▶ The keys  $x.key_i$  separate the ranges of keys stored in each subtree.

## Key ordering

- ▶ The keys  $x.key_i$  separate the ranges of keys stored in each subtree.
  - Let  $k_i$  be any key stored in the subtree with root  $x.c_i$

## Key ordering

- ▶ The keys  $x.key_i$  separate the ranges of keys stored in each subtree.
  - Let  $k_i$  be any key stored in the subtree with root  $x.c_i$ 
    - $k_i$  is one of the keys of child  $x.c_i$  of node  $x$

## Key ordering

- ▶ The keys  $x.key_i$  separate the ranges of keys stored in each subtree.
  - Let  $k_i$  be any key stored in the subtree with root  $x.c_i$ 
    - $k_i$  is one of the keys of child  $x.c_i$  of node  $x$
    - $k_i$  is one of the keys of a descendant of child  $x.c_i$  of node  $x$

## Key ordering

- ▶ The keys  $x.key_i$  separate the ranges of keys stored in each subtree.
  - Let  $k_i$  be any key stored in the subtree with root  $x.c_i$ 
    - $k_i$  is one of the keys of child  $x.c_i$  of node  $x$
    - $k_i$  is one of the keys of a descendant of child  $x.c_i$  of node  $x$

$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$




# Key ordering

- ▶ The keys  $x.key_i$  separate the ranges of keys stored in each subtree.

- Let  $k_i$  be any key stored in the subtree with root  $x.c_i$ 
  - $k_i$  is one of the keys of child  $x.c_i$  of node  $x$
  - $k_i$  is one of the keys of a descendant of child  $x.c_i$  of node  $x$

$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$

- ▶ Example:



## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?  $t$  children.

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?  $t$  children.
  - If the tree is nonempty, the root must have at least one key.



## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?  $t$  children.
  - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?  $t$  children.
  - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
  - Every node may contain at most  $2t - 1$  keys.

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?  $t$  children.
  - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
  - Every node may contain at most  $2t - 1$  keys.
  - So, how many children can an internal node have at most?

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?  $t$  children.
  - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
  - Every node may contain at most  $2t - 1$  keys.
  - So, how many children can an internal node have at most?  $2t$  children.

## Minimum degree

- ▶ Nodes have lower and upper bounds on the number of keys they can contain.
- ▶ Expressed in terms of a fixed integer  $t \geq 2$ , called the *minimum degree*
- ▶ Lower bound:
  - Every node other than the root must have at least  $t - 1$  keys.
  - So, how many children does an internal node need at least?  $t$  children.
  - If the tree is nonempty, the root must have at least one key.
- ▶ Upper bound:
  - Every node may contain at most  $2t - 1$  keys.
  - So, how many children can an internal node have at most?  $2t$  children.
  - We say that a node is *full* if it contains exactly  $2t - 1$  keys.

# Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory



## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:

## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
  - Fast.

## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
  - Fast.
  - Expensive and limited in capacity.

## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
  - Fast.
  - Expensive and limited in capacity.
- ▶ Secondary memory:

## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
  - Fast.
  - Expensive and limited in capacity.
- ▶ Secondary memory:
  - Cheaper and much larger.

## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
  - Fast.
  - Expensive and limited in capacity.
- ▶ Secondary memory:
  - Cheaper and much larger.
  - We need such capacity for storing large amounts of data, like in databases.


## Quick SOCS recap

- ▶ Computer systems use a hierarchy of memory technologies.

Registers < Caches < Main Memory < Secondary Memory

- ▶ Primary memory:
  - Fast.
  - Expensive and limited in capacity.
- ▶ Secondary memory:
  - Cheaper and much larger.
  - We need such capacity for storing large amounts of data, like in databases.
  - E.g. SSD's and HDD's.


## Secondary memory is slow – Example: HDD



- ▶ Because of the moving mechanical parts, HDD's are slow.
- ▶ SSD's are not mechanical, but still significantly slower than RAM.



## Secondary memory is slow – Example: HDD



- ▶ Because of the moving mechanical parts, HDD's are slow.
- ▶ SSD's are not mechanical, but still significantly slower than RAM.  
→ latency!

## Combating latency

- ▶ Access not just one item, but several at a time.

## Combating latency

- ▶ Access not just one item, but several at a time.
  - Information is divided into a number of equal-sized *blocks*.

## Combating latency

- ▶ Access not just one item, but several at a time.
  - Information is divided into a number of equal-sized *blocks*.
  - B-trees are a very common datastructure for handling this.

## Combating latency

- ▶ Access not just one item, but several at a time.
  - Information is divided into a number of equal-sized *blocks*.
  - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.

## Combating latency

- ▶ Access not just one item, but several at a time.
  - Information is divided into a number of equal-sized *blocks*.
  - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory  $\leftrightarrow$  reading/writing in a B-tree

# Combating latency

- ▶ Access not just one item, but several at a time.
  - Information is divided into a number of equal-sized *blocks*.
  - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory  $\leftrightarrow$  reading/writing in a B-tree
  - The number of blocks read or written provides a good approximation of the total time spent accessing the disk drive.

# Combating latency

- ▶ Access not just one item, but several at a time.
  - Information is divided into a number of equal-sized *blocks*.
  - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory  $\leftrightarrow$  reading/writing in a B-tree
  - The number of blocks read or written provides a good approximation of the total time spent accessing the disk drive.
  - Only downward paths.



## Combating latency

- ▶ Access not just one item, but several at a time.
  - Information is divided into a number of equal-sized *blocks*.
  - B-trees are a very common datastructure for handling this.
- ▶ Each node is usually as large as a whole disk block.
- ▶ Reading/writing in secondary memory  $\leftrightarrow$  reading/writing in a B-tree
  - The number of blocks read or written provides a good approximation of the total time spent accessing the disk drive.
  - Only downward paths.
  - So we want the height to be as small as possible.

# Outline

- 1 Definition of a B-tree
- 2 Use-cases
- 3 Height of a B-tree**
- 4 Operations on B-trees

# Height of a B-tree

## Theorem

*If  $n \geq 1$ , then for any  $n$ -key B-tree  $T$  of height  $h$  and minimum degree  $t \geq 2$*


$$h \leq \log_t \left( \frac{n+1}{2} \right)$$

# Height of a B-tree

## Theorem

If  $n \geq 1$ , then for any  $n$ -key B-tree  $T$  of height  $h$  and minimum degree  $t \geq 2$

$$h \leq \log_t \left( \frac{n+1}{2} \right)$$




# Height of a B-tree

## Theorem

If  $n \geq 1$ , then for any  $n$ -key B-tree  $T$  of height  $h$  and minimum degree  $t \geq 2$


$$h \leq \log_t \left( \frac{n+1}{2} \right)$$



depth	0	1	2	3	...	$h$
number of nodes	1	2	$2t$	$2t^2$	...	$2^{h-1}$

# Height of a B-tree

*Proof.*




depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
$\vdots$	$\vdots$
$h$	$2^{h-1}$

# Height of a B-tree

*Proof.* The number of keys  $n$  satisfies

$$n \geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}}$$




depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
$\vdots$	$\vdots$
$h$	$2^{h-1}$

# Height of a B-tree

*Proof.* The number of keys  $n$  satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left( \frac{t^h - 1}{t - 1} \right)
 \end{aligned}$$




depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
$\vdots$	$\vdots$
$h$	$2^{h-1}$



# Height of a B-tree

*Proof.* The number of keys  $n$  satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left( \frac{t^h - 1}{t-1} \right) \\
 &= 2t^h - 1
 \end{aligned}$$



depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
$\vdots$	$\vdots$
$h$	$2t^{h-1}$


# Height of a B-tree

*Proof.* The number of keys  $n$  satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2t^{i-1}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left( \frac{t^h - 1}{t - 1} \right) \\
 &= 2t^h - 1
 \end{aligned}$$

Thus

$$t^h \leq \frac{n+1}{2}$$



depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
$\vdots$	$\vdots$
$h$	$2t^{h-1}$


# Height of a B-tree

*Proof.* The number of keys  $n$  satisfies

$$\begin{aligned}
 n &\geq \underbrace{1}_{\text{root}} + \underbrace{(t-1)}_{\text{other nodes}} \sum_{i=1}^h \underbrace{2^{t^{i-1}}}_{\text{\#nodes}} \\
 &= 1 + 2(t-1) \left( \frac{t^h - 1}{t - 1} \right) \\
 &= 2t^h - 1
 \end{aligned}$$

Thus

$$t^h \leq \frac{n+1}{2} \iff h \leq \log_t \left( \frac{n+1}{2} \right)$$



depth	#nodes
0	1
1	2
2	$2t$
3	$2t^2$
$\vdots$	$\vdots$
$h$	$2^{h-1}$

## Comparison with red-black trees

- ▶ Height of the tree grows as  $O(\lg n)$  in both cases.

## Comparison with red-black trees

- ▶ Height of the tree grows as  $O(\lg n)$  in both cases.
  - Constant  $t$  is left out in asymptotic notation.

## Comparison with red-black trees

- ▶ Height of the tree grows as  $O(\lg n)$  in both cases.
  - Constant  $t$  is left out in asymptotic notation.
  - The base logarithm for B-trees is much larger:  $t$  vs 2.

## Comparison with red-black trees

- ▶ Height of the tree grows as  $O(\lg n)$  in both cases.
  - Constant  $t$  is left out in asymptotic notation.
  - The base logarithm for B-trees is much larger:  $t$  vs 2.
- ▶ Factor of  $\approx \lg t$  saved over red-black trees.

## Comparison with red-black trees

- ▶ Height of the tree grows as  $O(\lg n)$  in both cases.
  - Constant  $t$  is left out in asymptotic notation.
  - The base logarithm for B-trees is much larger:  $t$  vs 2.
- ▶ Factor of  $\approx \lg t$  saved over red-black trees.
- ▶ B-trees are limited in depth but contain an enourmous amount of keys.



## Comparison with red-black trees

- ▶ Height of the tree grows as  $O(\lg n)$  in both cases.
  - Constant  $t$  is left out in asymptotic notation.
  - The base logarithm for B-trees is much larger:  $t$  vs 2.
- ▶ Factor of  $\approx \lg t$  saved over red-black trees.
- ▶ B-trees are limited in depth but contain an enourmous amount of keys.
  - Few memory operations, so minimal memory latency.

# Outline

- ① Definition of a B-tree
- ② Use-cases
- ③ Height of a B-tree
- ④ Operations on B-trees

## Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.

## Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.

## Searching a B-tree


- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.
  - The amount of *ways* is equal to the number of children.

## Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.
  - The amount of *ways* is equal to the number of children.
  - At each internal node  $x$ , the search makes an  $(x.n + 1)$ -way branching decision.

## Searching a B-tree

- ▶ B-TREE-SEARCH is much like searching a binary tree.
- ▶ However, *multiway* (as opposed to binary) branching decisions are made.
  - The amount of *ways* is equal to the number of children.
  - At each internal node  $x$ , the search makes an  $(x.n + 1)$ -way branching decision.
- ▶ Example: searching for the letter  $S$ .



## Searching a B-tree: pseudocode

B-TREE-SEARCH( $x, k$ )

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```



## Searching a B-tree: pseudocode

B-TREE-SEARCH( $x, k$ )

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

►  $x$  is a pointer to the root node of a subtree

## Searching a B-tree: pseudocode

B-TREE-SEARCH( $x, k$ )

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- ▶  $x$  is a pointer to the root node of a subtree
- ▶  $k$  is the key to be searched

## Searching a B-tree: pseudocode

B-TREE-SEARCH( $x, k$ )

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- ▶  $x$  is a pointer to the root node of a subtree
- ▶  $k$  is the key to be searched
- ▶ If  $k$  is in the B-tree, B-TREE-SEARCH returns the node-index pair  $(y, i)$  such that  $y.key_i = k$ . Otherwise, the procedure returns NIL.

## Searching a B-tree: pseudocode

B-TREE-SEARCH( $x, k$ )

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- ▶  $x$  is a pointer to the root node of a subtree
- ▶  $k$  is the key to be searched
- ▶ If  $k$  is in the B-tree, B-TREE-SEARCH returns the node-index pair  $(y, i)$  such that  $y.key_i = k$ . Otherwise, the procedure returns NIL.

- ▶ Recall that  $k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$  where  $k_i$  is any key of child  $x.c_i$

## Searching a B-tree: memory accesses

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.

## Searching a B-tree: memory accesses

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.
- ▶ Thus, B-TREE-SEARCH executes  $O(h) = O(\log_t n)$  memory accesses.

## Searching a B-tree: memory accesses

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.
- ▶ Thus, B-TREE-SEARCH executes  $O(h) = O(\log_t n)$  memory accesses.
  - Recall that  $h \leq \log_t \left(\frac{n+1}{2}\right)$ .

## Searching a B-tree: memory accesses

- ▶ The nodes encountered during the recursion form a simple path downward from the root to the tree.
- ▶ Thus, B-TREE-SEARCH executes  $O(h) = O(\log_t n)$  memory accesses.
  - Recall that  $h \leq \log_t \left(\frac{n+1}{2}\right)$ .
  - Constants are left out in  $O$ -notation.



## Searching a B-tree: runtime

B-TREE-SEARCH( $x, k$ )

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- For every node  $x$ ,  $x.n < 2t$ , so the while loop takes  $O(t)$  time for each encountered node.

## Searching a B-tree: runtime

B-TREE-SEARCH( $x, k$ )

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

- ▶ For every node  $x$ ,  $x.n < 2t$ , so the while loop takes  $O(t)$  time for each encountered node.
- ▶ The total runtime is  $O(th) = O(t \log_t n)$ .

## Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.

## Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.
- ▶ We cannot simply create a new leaf node and insert it, because the resulting tree would fail to be a valid B-tree.


## Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.
- ▶ We cannot simply create a new leaf node and insert it, because the resulting tree would fail to be a valid B-tree.
- ▶ So, we need to insert the key into an *existing* leaf node.

## Inserting a key

- ▶ Significantly more complicated than inserting a key into a binary tree.
- ▶ We cannot simply create a new leaf node and insert it, because the resulting tree would fail to be a valid B-tree.
- ▶ So, we need to insert the key into an *existing* leaf node.
- ▶ What if a node  $y$  is *full*, i.e. what if  $y.n = 2t - 1$ ?

## Splitting a full node $y$ with $y.n = 2t - 1$



### B-TREE-SPLIT-CHILD( $x, i$ )

- ▶  $x$  is a non-full node and  $y = x.c_i$  is a full child of  $x$ .
- ▶ Split  $y$  about its median key  $S$  and move  $S$  up into  $y$ 's parent node  $x$ .
- ▶ Every key  $y.key_i$  that is greater than the median  $S$ , is placed in a new node  $z$ , which is a new child of  $x$ .

## Inserting a key (continued)

Complexity of B-TREE-INSERT( $T, k$ )?



## Inserting a key (continued)

Complexity of B-TREE-INSERT( $T, k$ )?

- ▶ Again  $O(h) = O(\log_t n)$  disk accesses.

## Inserting a key (continued)

Complexity of B-TREE-INSERT( $T, k$ )?

- ▶ Again  $O(h) = O(\log_t n)$  disk accesses.
- ▶ Again  $O(th) = O(t \log_t n)$  time required.

Questions?